

Computer System Organization

Recitation

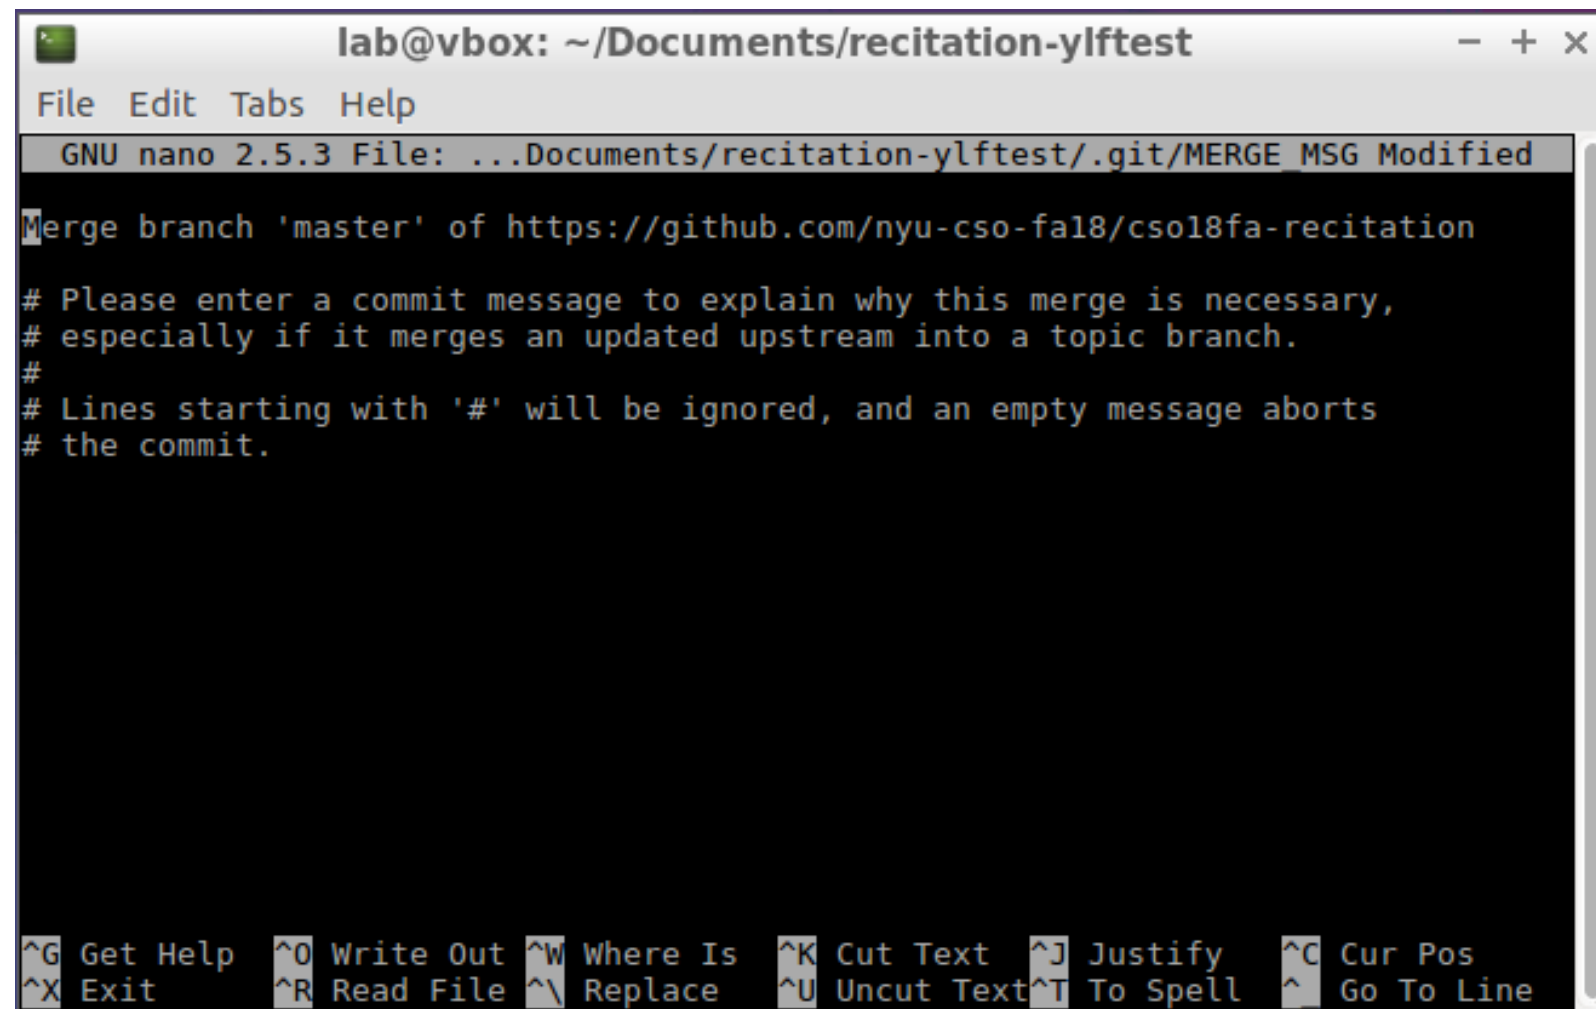
[Spring 2019]

CSCI-UA 0201-002

R3: Integer, Float, C Basics

Pull r03 from upstream

- git pull upstream master
- Get a merge message? Hit **ctrl + x**



The screenshot shows a terminal window titled 'lab@vbox: ~/Documents/recitation-ylftest'. The window contains the GNU nano 2.5.3 text editor. The file being edited is '...Documents/recitation-ylftest/.git/MERGE_MSG Modified'. The content of the file is a merge message template. The bottom status bar of the nano editor shows various keyboard shortcuts.

```
lab@vbox: ~/Documents/recitation-ylftest
File Edit Tabs Help
GNU nano 2.5.3 File: ...Documents/recitation-ylftest/.git/MERGE_MSG Modified
Merge branch 'master' of https://github.com/nyu-cso-fa18/cso18fa-recitation
# Please enter a commit message to explain why this merge is necessary,
# especially if it merges an updated upstream into a topic branch.
#
# Lines starting with '#' will be ignored, and an empty message aborts
# the commit.

^G Get Help  ^O Write Out ^W Where Is  ^K Cut Text  ^J Justify   ^C Cur Pos
^X Exit      ^R Read File ^\ Replace   ^U Uncut Text ^T To Spell  ^_ Go To Line
```

Starting from today

- We will do a quick review of lecture contents
 - Quick in-class questions
- Do recitation exercise together
 - You will have several minutes to read and do questions by yourself
 - After that I will release the answer in class

Today's topic


- Memory model
- How to represent integer
- How to represent real numbers
- Bit operators

Bits and Bytes

- Bit
 - smallest storage unit
 - two possible states (0, 1)
 - A typical “8GB” computer memory has 2^{36} bits
- Byte
 - 1 byte = 8 bits
 - 1 byte can represent $2^8 = 256$ possible states
 - computer memory is addressed at byte granularity

Memory model

- Everything is stored in memory
- Memory is a sequence of cells (logically)
- Each cell has 8 bits (1 byte)
- Each cell has an address (0-indexed)
 - What's address space of 16MB memory?
 - From 0 to $2^{24}-1$
- Think about everything from memory perspective



Things to memorize

- 1 byte = 8 bits
- common powers of 2: $2^0 = 1$, ..., $2^9 = 512$, $2^{10} = 1024$
- $2^{10} \text{ B} = 1 \text{ KB}$
- $2^{20} \text{ B} = 1 \text{ MB}$
- $2^{30} \text{ B} = 1 \text{ GB}$
- 2^{10} is approximately equal to 1000 (10^3)

Quick question

- Which of the following is closest to 1 million?
 - A. 2^{10}
 - B. 2^{20}
 - C. 2^{30}

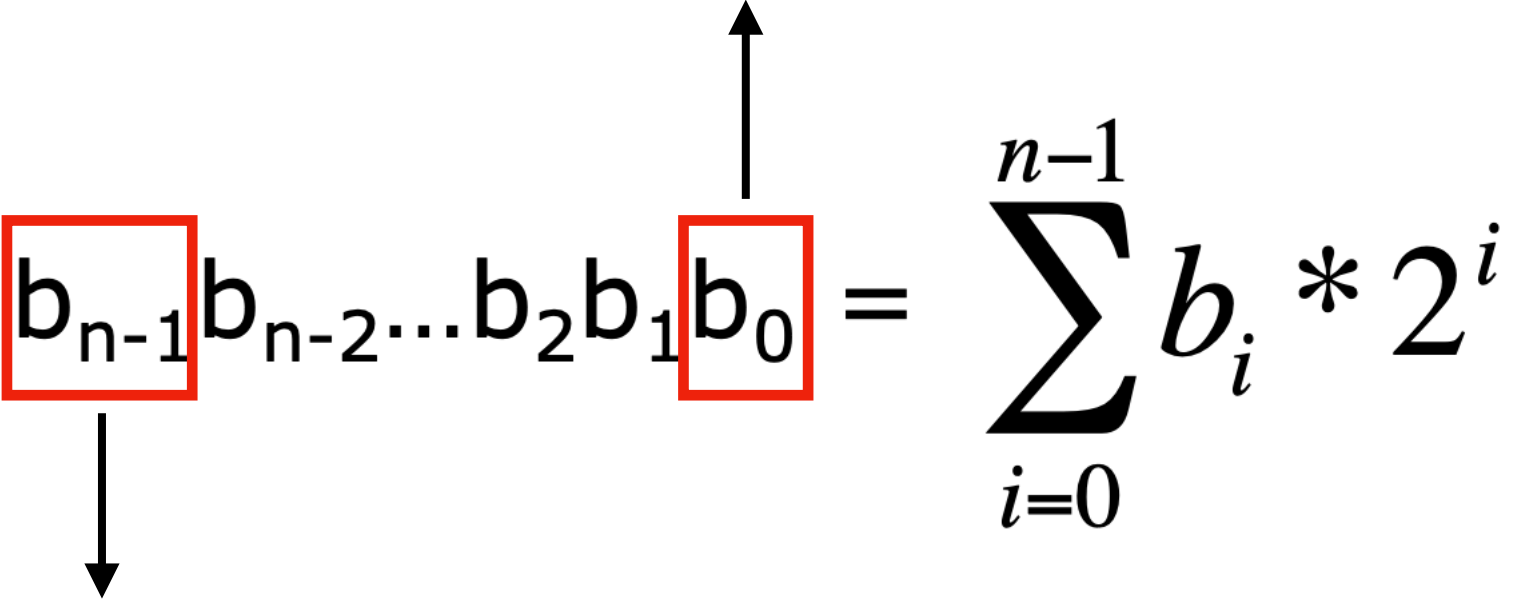
Representing Non-negative integer

- Base-2 representation

Least Significant Bit (LSB)

$b_{n-1}b_{n-2}\dots b_2b_1b_0 = \sum_{i=0}^{n-1} b_i * 2^i$

Most Significant Bit (MSB)



Exercise

- What's the decimal value of the binary representations?

Binary	Decimal	Hexadecimal
0000 0100	$2^2 = 4$	0x04
0010 0010	$2^5 + 2^1 = 34$	0x22
0101 1101	$2^6 + 2^4 + 2^3 + 2^2 + 2^0 = 93$	0x5D

Representing signed integers

- 2's complement

$$\vec{b} = [b_w, b_{w-1}, \dots, b_0] \qquad \text{val}(\vec{b}) = -b_w 2^w + \sum_{i=0}^{w-1} b_i 2^i$$

- How to tell is whether b is a negative number?
 - b is negative iff MSB b_w is 1

From 2's complement to decimal

- Method 1:

- Calculate $-b_w 2^w + \sum_{i=0}^{w-1} b_i 2^i$

Example:

1101 1000

$= -2^7 + 2^6 + 2^4 + 2^3$

$= -40$

- Method 2:

- rewrite the above formula

$$\begin{aligned} -b_w 2^w + \sum_{i=0}^{w-1} b_i 2^i &= -2b_w 2^w + b_w 2^w + \sum_{i=0}^{w-1} b_i 2^i \\ &= \sum_{i=0}^w b_i 2^i - b_w 2^{w+1} = \sum_{i=0}^w b_i 2^i - b_w 2^n \end{aligned}$$

Example:

1101 1000

$= 2^7 + 2^6 + 2^4 + 2^3 - 2^8$

$= 216 - 256 = -40$

- So, view the binary as non-negative number, and then subtract $b_w 2^n$, n is number of bits

From negative decimal to 2's complement

- Method 1

- Convert the absolute value to binary format
- flip all bits
- add 1 (don't forget me !!!)

Example:
-40

0010 1000

1101 0111

1101 1000

- Method 2

- Remember decimal value $D = \sum_{i=0}^w b_i 2^i - b_w 2^n$
- Then $\sum_{i=0}^w b_i 2^i = D + b_w 2^n$
- Add D by 2^n and convert to binary, n is number of bits

Example
-40 + 2^8 = 216
= 1101 1000

Integer types in C

- You need to memorize the size!

64 bits machine

type	size (bytes)	example
(unsigned) char	1	char c = 'a'
(unsigned) short	2	short s = 12
(unsigned) int	4	int i = 1
(unsigned) long	8	long l = 1

Exercise

- What's the decimal value of the following numbers?

Hex	Binary	char	unsigned char	short	unsigned short
0xB6	1011 0110	$-2^7+2^5+2^4+2^2+2^1=-74$	$2^7+2^5+2^4+2^2+2^1=182$	182	182
0x7A	0111 1010	$2^6+2^5+2^4+2^3+2^1=122$	122	122	122

Exercise




- Do it again faster
 - Convert from hex to decimal directly!

Hex	char	unsigned char	short	unsigned short
0xB6	$11 \cdot 16 + 6 - 2^8$ =-74	$11 \cdot 16 + 6 = 182$	182	182
0x7A	$7 \cdot 16 + 10 = 122$	122	122	122

From memory perspective

- How are multi-byte data stored in computer memory?
- Use C int type as an example:

int: 4-byte (32-bit)



Advantages of little-endian machine

- Overlap computation and memory read
 - Computation usually start from lower bits
 - simultaneously perform memory transfer and calculation

high	0000 0000
	0000 0000
	0100 0110
low	1000 0010

- Faster shortening
 - Casting from int to short?
 - ▶ Just read lower two bytes
 - What's the value if cast to char?
 - What's the value if cast to unsigned char?

**Unless otherwise noted,
we always assume 64-bit
little endian machine**

Integer overflow

- Assuming one-byte signed integer

$$\begin{array}{r} 1\ 0\ 0\ 0\ 0\ 0\ 0\ 1 \\ +\ 1\ 0\ 0\ 0\ 0\ 0\ 0\ 1 \\ \hline 1\ 0\ 0\ 0\ 0\ 0\ 0\ 1\ 0 \end{array}$$

carry bit discarded

2 ???

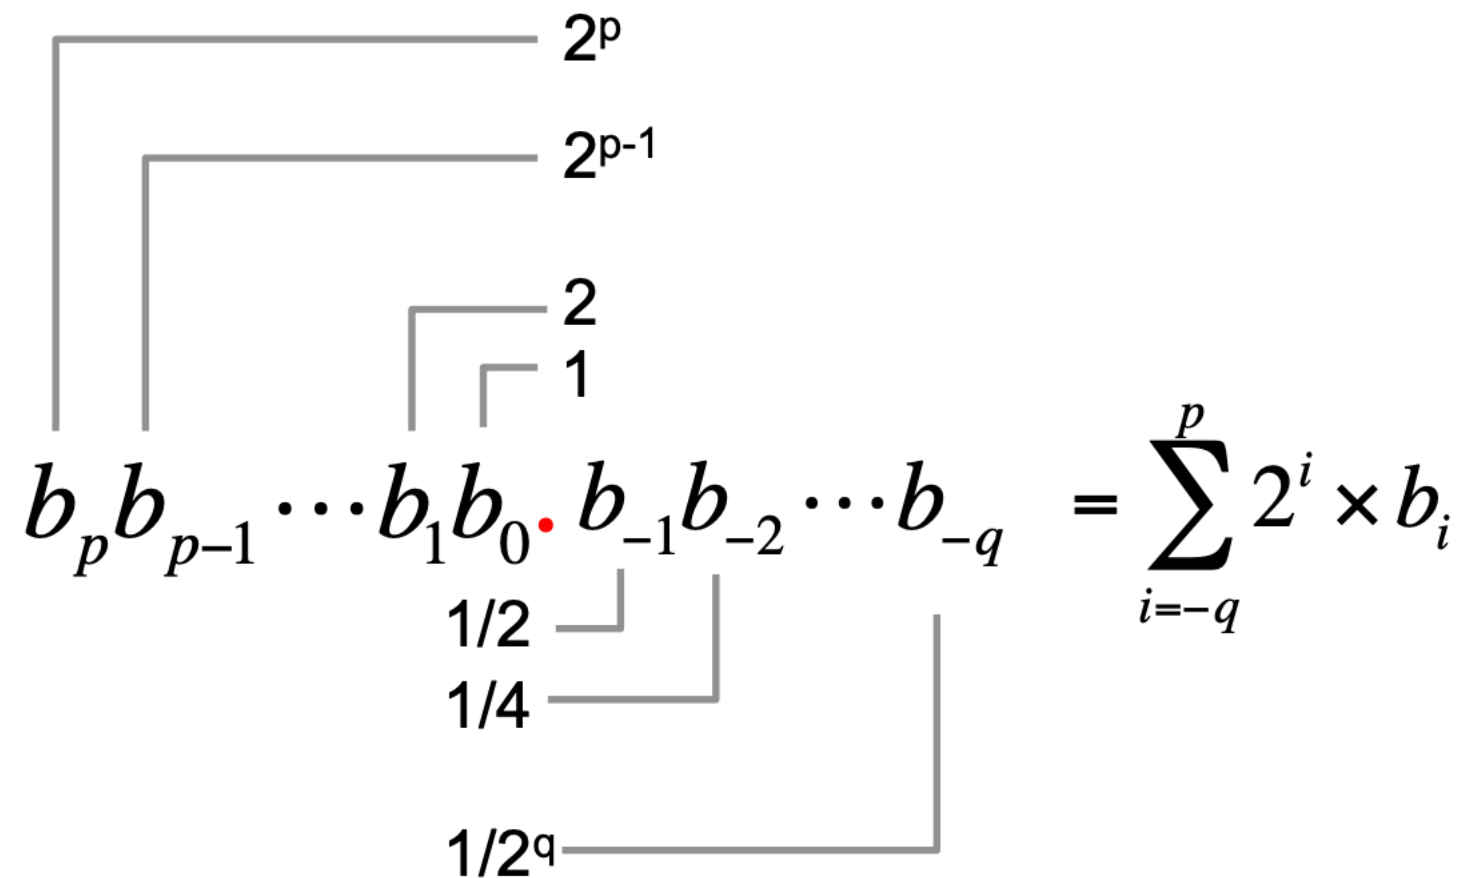
- How to tell if there is an integer overflow
 - pay attention to the data type
 - whenever the computation result goes out of range, it's an overflow
 - for example, $1 + 127 = -128$; $-128 - 1 = 127$; ...

Representing real numbers

Converting real numbers from decimal to binary

$$r_{10} = (d_m d_{m-1} d_1 d_0 \cdot d_{-1} d_{-2} \dots d_{-n})_{10}$$

$$= (b_p b_{p-1} b_1 b_0 \cdot b_{-1} b_{-2} \dots b_{-q})_2$$




$$b_p b_{p-1} \dots b_1 b_0 \cdot b_{-1} b_{-2} \dots b_{-q} = \sum_{i=-q}^p 2^i \times b_i$$

Single-precision floating point number

$$r_{10} = \pm M * 2^E, \text{ where } 1 \leq M < 2$$

$$M = (1.b_1b_2b_3...b_{23})_2$$

M: significant, E: exponent



$$(b_1b_2b_3...b_{23})_2$$

Floating bias 127 for single-precision

Why 127?

**Pay attention
to the order**

Exercise

- If the right-hand side figure is the memory representation of a float number, what's its decimal value?

high	0x3F
	0x00
	0x00
low	0x00

Little-endian machine (by default)

$0x3F000000 = (0011\ 1111\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000)_2$

$$S = 0$$

$$E = (01111110)_2 - 127 = 126 - 127 = -1$$

$$F = 0$$

$$M = 1.0$$

$$\text{Decimal value} = +1.0 * 2^{-1} = 0.5$$

Exercise

- What's hex representation of float number -5.5625?

$$(-5.5625)_{10} = (-101.1001)_2 = (-1.011001 * 2^2)_2$$

$$S = 1$$

$$E = 2 + 127 = 129 = (1000\ 0001)_2$$

$$M = 1.011001$$

$$F = 011001$$


$$(1100\ 0000\ 1011\ 0010\ 0000\ 0000\ 0000\ 0000)_2$$

$$= 0xC0B20000$$

IEEE denormalized representation


- Why?
 - Represent 0 and values very close to 0
- How?

Denormalized Encoding:




Zeros

+0




-0



$$+0 = -0$$

Special Values

Special Value's Encoding:



values	sign	frac	Computation Rules
$+\infty$	0	all zeros	$1/+0 = +\infty$, $3e38+1e38 = +\infty$, $1/\infty = +0$
$-\infty$	1	all zeros	$1/-0 = -\infty$, $-3e38-1e38 = -\infty$, $1/-\infty = -0$
NaN	any	non-zero	$\text{sqrt}(-1)$, $\infty + \infty$

Other things you need to know

- Double-precision format
 - S (1 bit), E (11 bits), F(52 bits)
- Using floating points are tricky
 - Invalid operation: $0/0$, $\text{sqrt}(-1)$, $\infty + \infty$
 - Divide by zero: $x/0 = \infty$
 - Overflows: result too big to fit
 - Underflows: $0 < \text{result} < \text{smallest denormalized value}$
 - Inexact: rounding!

C bit operators

C operators

- Similar to other languages like Java and Python
 - operations may have different symbols in different language

Arithmetic

`+, -, *, /, %, ++, --`

Relational

`==, !=, >, <, >=, <=`

Logical

`&&, ||, !`

Bitwise

`&, |, ^, ~, >>, <<`

Bit operators

- Operators that perform bit-wise operations:
 - `&`: bit-wise and
 - `|`: bit-wise or
 - `^`: bit-wise xor
 - `~`: bit-wise negate (flip bits)
 - `>>`, `<<`: right and left shift bits
- Why bit operators?
 - Special functionalities
 - Fast! Can usually be done in parallel

How to write multiply by 2?

- Use multiply operator
 - $x = x * 2;$
 - Complexity: $O(n^2)$, n is number of bits
- Use add operator
 - $x = x + x;$
 - Complexity: $O(n)$
- Use bit operator
 - $x = x \ll 1;$
 - Complexity: $O(1)$

Shift operator

- \ll shift all bits left, and pad 0 at right side
 - \gg shift bits right
 - Logical shift: always pad 0 at left side
 - Arithmetic shift: pad sign bit (MSB) at left side
 - Why do we need arithmetic shift?
 - Semantically, shift means multiply or divide by some power of 2
 - If pad 0, $(-2)_{10} \gg 1 = (1111\ 1110)_2 \gg 1 = (0111\ 1111)_2 = (127)_{10}$
 - Which shift does C use?
 - Unsigned type uses logical shift
 - Signed type uses arithmetic shift
- Why can't unsigned type use arithmetic shift as well?**

Operator precedence

- What's the result of the following expression?
 - `x = 16 >> 3 + 1;`
- bit operator has lower precedence than arithmetic operator
 - Use parenthesis whenever you use bit operators!

Lab1

- Mini part:
 - 8 simple C exercises
 - complete and test each exercise individually
 - don't work in a failure-stop mode!
- Scratch part
 - write stand-alone program from scratch
 - write your own Makefile
 - learn how to use IO to read write files
- Pay attention to coding style (20% of your lab grade)
 - <http://cs61.seas.harvard.edu/wiki/2015/Style>

Lab1

- Read (updated) lab1 instructions carefully
 - <https://nyu-cso.github.io/labs/l1.html>
- Lab 1 due 2/25 11pm
- Make sure your code is tested inside VM
- Pay attention to the rules:
 - only modify files you are allow to change
 - check your scratch part output format using the `output_format_checker.sh`
- Get crazy pointer bugs and segmentation fault?
 - Try GDB first!
 - Come to office hour, we will help you debug for lab 1

Today's exercise

- go through the readme and e1.c e2.c
- answer questions and complete code